

Final Project for Math 390 at Queens College
May 24, 2020

By [Steven Grgas]
In collaboration with:
[Andrew Claros]

Abstract

When buying or selling an apartment, what is the first piece of information that one seeks out? Often, people would like to know the listing price. Everyone lives on a budget and needs to know how much things cost. But the listing price is not as important as the sale price. We would like to know how accurate the listing price is to the sale price when we buy or sell an apartment. Websites like Zillow.com design algorithms that give us an estimate value of each home. There are many factors that go into determining a fair price. For instance, condos tend to be more expensive than co-ops. Zillow's algorithms know this because of the plentiful data that goes into them that comprises of many variables. However, if you consider Zillow's rate of success, it is easy to see that they are fairly inaccurate. In order to beat Zillow's estimate, we use many tricks available in the machine learning toolbox.

1. Introduction

The motivation for using algorithms is that we don't know the actual value of a home before it is sold. If we can get an estimate value of homes, we can help the seller and potential buyers influence their financial decision making. In order to provide an estimate value, we will need to use machine learning skills. First, we need a dataset; columns for our variables and rows for each home. What the algorithms will do is take in data and provide us with a function that will allow us to estimate the value of a new home, given some new data. This will be our response. The algorithms that were used were OLS, regression trees and random forests. The performance results were the best for random forests and worst for OLS. It was noticed that the data was negatively skewed, meaning that there were more homes with high costs. This would make the data points nonlinear, which means OLS is not an accurate model.

2. The Data

The data used for this project originally came from Amazon's MTurk. This provided us with a data frame of size 2231 x 55 of some relevant information on characteristics of specific apartments and the final sale price. It provides some important information as to how certain data could influence the sale price, including certain features of the homes such as number of rooms, square footage, etc. However, the model was underfit and the RSME was very large. The data was helpful as it represented what people valued in a home but it was not enough. We decided to add data from zipdatamaps.com. We supplemented our data with fifteen columns of relevant data, making our new data frame have size 2231 x 70. This new data consisted of sociological data that are directly correlated with home prices. There were some outliers in the sale

price for rich and poor neighborhoods. Extrapolation of the model is likely with some of the features. The sociological features such as unemployment rate, median household income, population density and race are destined to change over time. Renovation of apartments may change the number of floors, number of rooms, and taxes. This data should be kept up to date as one continues to run the model for future costs of apartments.

2.2 Featurization

The data provided by MTurk and zipdatamaps.com was direct for this project and we did not need to featurize any data provided.

2.3 Errors and Missingness

The data frame from MTurk occasionally contained misspelled words, '0' or '1' instead of 'yes' or 'no,' and missing data. To correct the mistypes, we simply changed the misspelled words to their correct spelling and turned a '0' to 'no' and '1' to 'yes' because we needed consistency in order to run the algorithms. When viewing the data we observed that many columns were missing essential data. If the column was at least 50% complete, we would impute the missing data via the random forest algorithm. This allows us to enjoy running algorithms that contain quality features for our experiment without computing nonsense results.

3.1 Regression Tree Modeling

The top ten features for regression trees were number of full bathrooms, apartment type, approximate year built, common charges, number of bedrooms, number of bathrooms, maintenance cost, number of floors in building, walk score, and kitchen type combo.

3.2 Linear Modeling

The OLS algorithm had in-sample RMSE = \$71,000 and $R^2 = 0.83$. In-sample RMSE is high, but this is expected since the data is nonlinear. R^2 scored better than expected, but the model is not accurate enough to ship out. The ten best features for regression tree have the following coefficients for the linear model: Number of full bathrooms = 8.4×10^4 , apartment type = 2.46×10^5 , approximate year built = -2.27×10^5 , common charges = 1.4×10 , number of bedrooms = 5×10^4 , maintenance cost = 1.16×10^2 , number of floors in building = 2.215×10^3 , walk score = -1.8×10^2 , kitchen type(eat-in) = -1.45×10^4 , kitchen type(efficiency) = -2.02×10^4 . If z represents the coefficient for one of these variables, then a change in one unit for the x value will result in a shift by z .

3.3 Random Forest Modeling

Since OLS is a linear model, it can be ruled out as the best model. Regression trees is a good algorithm, but random forests always beats it because it trains on subsets of the data and chooses the best one. Regression trees work on one data set. The oos RMSE = \$74,403 for regression trees.

4. Performance Results for Random Forest Modeling

$R^2 = 0.908$ which means the model was very accurate. RMSE = \$52,795. This was the lowest RMSE out of each algorithm. For generalization error, we would estimate $R^2 = 0.85$ and RMSE = \$60,000 due to change in many factors of real estate.

5. Discussion

The biggest problem with this project was the lack of good data. There was a lot of missingness that had to be imputed. Once imputed, the errors were still high, so we needed to seek elsewhere for good data. This is when we thought about social issues such

as population density and income. These factors largely influence housing costs so we felt they were important to include. In order to keep this model accurate in time, relevant information needs to be updated.

Acknowledgments

Family member Dan Price who helped with his knowledge of real state.

References

zipdatamaps.com

Data Munging

```
pacman::p_load(dplyr, skimr, lubridate, stringr, tidyr, ggplot2, randomForest, caTools, missForest, mice, VIM, YARF, mlr, data.table)
```

```
## YARF can now make use of 7 cores.
```

```
require(caTools)
```

```
scraped = read.csv('scraped_demographics.csv')
```

```
scraped=scraped%>%
```

```
  select(-c(X, racial_majority, url, population, total_households, inc_below50, inc_bet_50_100, inc_above_100, average_sale_price))%>%
```

```
  mutate(zipcode = as.factor(zipcode))
```

```
scraped
```

##	zipcode	town	med_income	pop_density
## 1	11361	bayside	74579	2.687271
## 2	11362	little	79329	2.538889
## 3	11363	little	89073	2.618209
## 4	11364	oakland	73011	2.600662
## 5	11354	flushing	47974	2.727535
## 6	11355	flushing	48502	2.977806
## 7	11356	college	63255	3.023868
## 8	11357	whitestone	73016	2.584329
## 9	11358	flushing	68523	2.806129
## 10	11359	bayside	78919	NA
## 11	11360	bayside	77626	2.156200
## 12	11365	fresh	69995	2.732458
## 13	11366	fresh	82431	2.994468
## 14	11367	flushing	55313	2.741584
## 15	11412	saint	62628	3.262744
## 16	11423	hollis	60339	3.137044
## 17	11432	jamaica	56699	3.296238
## 18	11433	jamaica	41093	3.195210
## 19	11434	jamaica	56783	2.920816
## 20	11435	jamaica	51678	3.029057
## 21	11436	jamaica	51051	3.353073
## 22	11101	long	36133	2.443571
## 23	11102	astoria	45715	2.423702
## 24	11103	astoria	51217	2.296033
## 25	11104	sunnyside	49244	2.301750
## 26	11105	astoria	50741	2.275931
## 27	11106	astoria	45208	2.273259
## 28	11374	rego	53555	2.257781
## 29	11375	forest	69665	2.111808
## 30	11379	middle	64453	2.508537
## 31	11385	ridgewood	47716	2.829607

```
## 32 11004      glen      74878    2.817286
## 33 11005      floral    146600    1.325991
## 34 11411      cambria    80416    3.148821
## 35 11413 springfield    72995    3.182465
## 36 11422      rosdale    76463    3.233950
## 37 11426      bellerose  75884    2.904557
## 38 11427      queens    66639    3.070806
## 39 11428      queens    71446    3.459928
## 40 11429      queens    42750    3.443759
## 41 11414      howard    67161    2.444881
## 42 11415      kew       60876    2.245298
## 43 11416      ozone     51482    3.440017
## 44 11417      ozone     56312    3.138353
## 45 11418      richmond   53300    3.286737
## 46 11419      south     55072    3.720917
## 47 11420      south     61621    3.433504
## 48 11421      woodhaven  58075    3.303250
## 49 11368      corona     45741    3.808981
## 50 11369      east      51217    3.449616
## 51 11370      east      44700    4.215401
## 52 11372      jackson    50985    2.787417
## 53 11373      elmhurst   48378    3.216154
## 54 11377      woodside   49306    2.859735
## 55 11378      maspeth    56961    2.741028
```

```
housing_orig = read.csv("housing_data_2016_2017.csv")
set.seed(10000)
```

```
housing= housing_orig[,29:ncol(housing_orig)]
```

```
housing= housing %>%
  mutate(common_charges =as.numeric(gsub('[$, ]','',common_charges))) %>%
  mutate(dogs_allowed= factor(replace(dogs_allowed, str_detect(tolower(dogs_a
llowed),pattern='y'),'yes')) %>%
  mutate(cats_allowed= factor(replace(cats_allowed, str_detect(tolower(cats_a
llowed),pattern='y'),'yes')) %>%
  mutate(fuel_type= na_if(fuel_type,c('Other')) %>%
  mutate(fuel_type= na_if(fuel_type,c('other')) %>%
  mutate(fuel_type= replace(fuel_type, str_detect(fuel_type,pattern='oil'),'g
as')) %>%
  mutate(fuel_type= replace(fuel_type, str_detect(fuel_type,pattern='none'),N
A)) %>%
  mutate(fuel_type= factor(fuel_type)) %>%
  mutate(kitchen_type= na_if(kitchen_type,c('1955')) %>%
  mutate(kitchen_type= na_if(kitchen_type,c('none')) %>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type,pattern=
'Eat in'),'eatin')) %>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type,pattern=
```

```

'Eat In'), 'eatin'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'eat in'), 'eatin'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'Combo'), 'combo'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'efficiemcy'), 'efficiency'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'efficiency kitchen'), 'efficiency'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'efficiency kitchen'), 'efficiency'))%>%
  mutate(kitchen_type= replace(kitchen_type, str_detect(kitchen_type, pattern=
'efficiency ktchen'), 'efficiency'))%>%
  mutate(kitchen_type= factor(kitchen_type))%>%
  mutate(maintenance_cost =as.numeric(gsub('[$, ]', '', maintenance_cost)))%>%
  mutate(parking_charges =as.numeric(gsub('[$, ]', '', parking_charges)))%>%
  mutate(sale_price =as.numeric(gsub('[$, ]', '', sale_price)))%>%
  mutate(total_taxes =as.numeric(gsub('[$, ]', '', total_taxes)))%>%
  mutate(listing_price_to_nearest_1000 = as.numeric(gsub('[$, ]', '', listing_p
rice_to_nearest_1000)))%>%
  mutate(zipcode = str_extract(full_address_or_zip_code, '[0-9]{5}'))

```

```
summary(housing$town)
```

```
## Length Class Mode
##      0  NULL  NULL
```

```
test_df = housing%>%
  drop_na(sale_price)
```

```
housing = housing %>%
  drop_na(sale_price) %>%
  drop_na(approx_year_built) %>%
  drop_na(kitchen_type)%>%
  drop_na(num_bedrooms)%>%
  drop_na(num_total_rooms)%>%
  drop_na(zipcode)
```

```
#to beat
```

```
housing = select(housing, -c('model_type', 'full_address_or_zip_code', 'url', 'g
arage_exists', 'pct_tax_deductibl', 'num_half_bathrooms', 'date_of_sale', 'listin
g_price_to_nearest_1000', 'community_district_num', 'dining_room_type'))
housing= inner_join(housing, scraped)
```

```
## Joining, by = "zipcode"
```



```
## Warning: Column `zipcode` joining character vector and factor, coercing in
to
## character vector
```

```
housing=select(housing,-('zipcode'))
housing=housing%>%
  group_by(town)%>%
  filter(n()>10)
```

```
#housing = dplyr::select(housing, -c('model_type','full_address_or_zipcode','
url','garage_exists','pct_tax_deductibl','num_half_bathrooms','date_of_sale',
'zipcode','listing_price_to_nearest_1000','community_district_num','dining_ro
om_type','fuel_type','sq_footage'))
```

```
y= housing$sale_price
sale_ind = which(colnames(housing)=='sale_price')
impute=mice(housing[, -sale_ind], seed = 10000)
```

```
##
## iter imp variable
## 1 1 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 1 2 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 1 3 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 1 4 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 1 5 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 2 1 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 2 2 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 2 3 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 2 4 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 2 5 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 3 1 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 3 2 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 3 3 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 3 4 common_charges fuel_type maintenance_cost num_floors_in_buildi
```

```

ng parking_charges sq_footage total_taxes
## 3 5 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 4 1 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 4 2 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 4 3 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 4 4 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 4 5 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 5 1 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 5 2 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 5 3 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 5 4 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes
## 5 5 common_charges fuel_type maintenance_cost num_floors_in_buildi
ng parking_charges sq_footage total_taxes

```

```
## Warning: Number of logged events: 175
```

```
colnames(housing)
```

```

## [1] "approx_year_built"      "cats_allowed"           "common_charges"
## [4] "coop_condo"            "dogs_allowed"           "fuel_type"
## [7] "kitchen_type"          "maintenance_cost"       "num_bedrooms"
## [10] "num_floors_in_building" "num_full_bathrooms"     "num_total_rooms"
## [13] "parking_charges"       "sale_price"             "sq_footage"
## [16] "total_taxes"           "walk_score"             "town"
## [19] "med_income"            "pop_density"

```

```
#sum(is.na(housing_cols_removed$total_taxes))/nrow(housing_cols_removed)
```

```

housing_imp=complete(impute,1)
housing_imp$sale_price = y
housing_df= housing_imp %>%
  drop_na(sale_price)
train_indices = sample(1 : nrow(housing_df), round((1 - .2) * nrow(housing_df
)))
housing_train = housing_df[train_indices, ]
test_indices=setdiff(1 : nrow(housing_df), train_indices)

housing_test= housing_df[test_indices,]

Y_test = housing_test$sale_price
X_test = housing_test

```

```

Y_train = housing_train$sale_price
X_train = housing_train
X_train$sale_price = NULL
n_train = nrow(X_train)

tree_mod=YARFCART(X_train, Y_train, calculate_oob_error = FALSE)

## YARF initializing with a fixed 1 trees...
## YARF factors created...
## YARF after data preprocessed... 38 total features...
## Beginning YARF regression model construction...done.

#get training performance
y_hat_train = predict(tree_mod, housing_train)

## Warning in predict.YARF(tree_mod, housing_train): Prediction set column na
mes did not match training set column names.
## Attempting to subset to training set columns.

e = Y_train - y_hat_train
tree_train_perf = 1 - sd(e) / sd(Y_train)
#test performance
y_hat_test = predict(tree_mod, housing_test)

## Warning in predict.YARF(tree_mod, housing_test): Prediction set column nam
es did not match training set column names.
## Attempting to subset to training set columns.

e = Y_test - y_hat_test
tree_test_perf = 1 - sd(e) / sd(Y_test)
#linear train
linear_mod = lm(sale_price ~ ., housing_train)
y_hat_train_linear = predict(linear_mod, housing_train)
e = Y_train - y_hat_train_linear
linear_train_perf = 1 - sd(e) / sd(Y_train)
#linear test
y_hat_test_linear = predict(linear_mod, housing_test)
e = Y_test - y_hat_test_linear
linear_test_perf = 1 - sd(e) / sd(Y_test)
SSE = sum(linear_mod$residuals^2)
MSE <- SSE / length(linear_mod$residuals)
RMSE <- sqrt(MSE)

paste0('tree in sample: ',tree_train_perf)

## [1] "tree in sample: 0.863321950183148"

paste0('tree oos: ',tree_test_perf)

## [1] "tree oos: 0.495599997678257"

```

```

paste0('linear in sample: ',linear_train_perf)

## [1] "linear in sample: 0.594198348862615"

paste0('linear oos: ',linear_test_perf)

## [1] "linear oos: 0.593888242765756"

paste0("lm RMSE: ", RMSE)

## [1] "lm RMSE: 71094.9678873884"

paste0("lm r.squared:", summary(linear_mod)$r.squared)

## [1] "lm r.squared:0.835325019934172"

summary(linear_mod)

##
## Call:
## lm(formula = sale_price ~ ., data = housing_train)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -300038  -40020   -1863   35142  307269
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -2.277e+05  7.521e+05  -0.303  0.762308
## approx_year_built    2.130e+02  3.739e+02   0.570  0.569363
## cats_allowedyes     1.476e+03  1.402e+04   0.105  0.916225
## common_charges     1.407e+01  2.976e+01   0.473  0.636763
## coop_condocondo     2.467e+05  2.011e+04  12.268 < 2e-16 ***
## dogs_allowedyes     1.302e+04  1.523e+04   0.855  0.393017
## fuel_typegas       7.538e+04  2.764e+04   2.727  0.006759 **
## kitchen_typeeatin   -1.456e+04  1.305e+04  -1.116  0.265368
## kitchen_typeefficiency -2.024e+04  1.251e+04  -1.617  0.106827
## maintenance_cost    1.159e+02  2.822e+01   4.107  5.16e-05 ***
## num_bedrooms       5.453e+04  1.002e+04   5.445  1.07e-07 ***
## num_floors_in_building 2.215e+03  8.739e+02   2.535  0.011740 *
## num_full_bathrooms   8.411e+04  1.459e+04   5.763  2.02e-08 ***
## num_total_rooms     8.456e+03  6.784e+03   1.247  0.213520
## parking_charges    -9.751e+01  8.826e+01  -1.105  0.270101
## sq_footage        -2.355e+01  1.486e+01  -1.585  0.114098
## total_taxes        -1.958e-01  3.027e+00  -0.065  0.948458
## walk_score        -1.826e+02  5.134e+02  -0.356  0.722255
## townbayside        -1.021e+05  7.155e+04  -1.427  0.154652
## towncorona        -2.675e+05  7.792e+04  -3.434  0.000678 ***
## townflushing      -1.274e+05  3.983e+04  -3.200  0.001519 **
## townforest        -2.737e+04  5.761e+04  -0.475  0.635096
## townglen         -7.091e+04  7.414e+04  -0.956  0.339598
## townhoward       -2.364e+05  5.707e+04  -4.142  4.45e-05 ***

```

```

## townjackson      -2.943e+04  4.209e+04  -0.699  0.484840
## townjamaica      -2.072e+05  5.489e+04  -3.776  0.000192 ***
## townkew         -1.544e+05  4.592e+04  -3.362  0.000873 ***
## townlittle      -9.024e+04  8.008e+04  -1.127  0.260668
## townoakland     -1.245e+05  6.728e+04  -1.851  0.065172 .
## townrego        -1.302e+05  4.111e+04  -3.166  0.001699 **
## townwhitestone  -1.224e+05  6.868e+04  -1.783  0.075626 .
## med_income      -1.978e+00  2.075e+00  -0.954  0.341017
## pop_density     -1.580e+04  4.710e+04  -0.336  0.737443
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 74830 on 306 degrees of freedom
## Multiple R-squared:  0.8353, Adjusted R-squared:  0.8181
## F-statistic: 48.51 on 32 and 306 DF,  p-value: < 2.2e-16

train_indices = sample(1 : nrow(housing_df), round((1 - .2) * nrow(housing_df)))
test_indices=setdiff(1 : nrow(housing_df), train_indices)
housing_train = housing_df[train_indices,]

y=housing_train %>%
  dplyr::select(sale_price)

rf <- makeLearner("regr.randomForest", predict.type = "response", par.vals =
list(ntree = 200, mtry = 3))
modeling_task_train = makeRegrTask(data = housing_train, target = "sale_price")

## Warning in makeTask(type = type, data = data, weights = weights, blocking =
## blocking, : Empty factor levels were dropped for columns: town

modeling_task_test = makeRegrTask(data = housing_test, target = "sale_price")

## Warning in makeTask(type = type, data = data, weights = weights, blocking =
## blocking, : Empty factor levels were dropped for columns: town

#instantiate the task
algorithm = makeLearner("regr.randomForest")
validation = makeResampleDesc("CV", iters = 5)
learner= makeLearner("regr.randomForest")
rf_param = makeParamSet(
  makeIntegerParam("ntree", lower = 50, upper = 500),
  makeIntegerParam("mtry", lower = 3, upper = 10),
  makeIntegerParam("nodesize", lower = 10, upper = 50)
)
rancontrol = makeTuneControlRandom(maxit = 50L)
set_cv = makeResampleDesc("CV", iters = 3L)

```

```
rf_tune = tuneParams(learner = algorithm, resampling = set_cv, task = modelin  
g_task_train, par.set = rf_param, control = rancontrol)
```

```
## [Tune] Started tuning learner regr.randomForest for parameter set:
```

```
##           Type len Def      Constr Req Tunable Trafo  
## ntree     integer - - 50 to 500 - TRUE -  
## mtry      integer - - 3 to 10 - TRUE -  
## nodesize integer - - 10 to 50 - TRUE -
```

```
## With control class: TuneControlRandom
```

```
## Imputation value: Inf
```

```
## [Tune-x] 1: ntree=128; mtry=9; nodesize=29
```

```
## [Tune-y] 1: mse.test.mean=7749229657.7481327; time: 0.0 min
```

```
## [Tune-x] 2: ntree=91; mtry=3; nodesize=12
```

```
## [Tune-y] 2: mse.test.mean=7600745431.9054537; time: 0.0 min
```

```
## [Tune-x] 3: ntree=337; mtry=9; nodesize=13
```

```
## [Tune-y] 3: mse.test.mean=6664777976.1915483; time: 0.0 min
```

```
## [Tune-x] 4: ntree=164; mtry=10; nodesize=25
```

```
## [Tune-y] 4: mse.test.mean=7506188461.1436682; time: 0.0 min
```

```
## [Tune-x] 5: ntree=216; mtry=4; nodesize=22
```

```
## [Tune-y] 5: mse.test.mean=8017497816.2785578; time: 0.0 min
```

```
## [Tune-x] 6: ntree=177; mtry=8; nodesize=23
```

```
## [Tune-y] 6: mse.test.mean=7520331450.3883829; time: 0.0 min
```

```
## [Tune-x] 7: ntree=429; mtry=8; nodesize=37
```

```
## [Tune-y] 7: mse.test.mean=8337884232.0532780; time: 0.0 min
```

```
## [Tune-x] 8: ntree=223; mtry=5; nodesize=39
```

```
## [Tune-y] 8: mse.test.mean=8881813980.1585846; time: 0.0 min
```

```
## [Tune-x] 9: ntree=318; mtry=5; nodesize=40
```

```
## [Tune-y] 9: mse.test.mean=8743730708.1664810; time: 0.0 min
```

```
## [Tune-x] 10: ntree=465; mtry=3; nodesize=49
```

```
## [Tune-y] 10: mse.test.mean=10231063501.3951530; time: 0.0 min
```

```
## [Tune-x] 11: ntree=310; mtry=8; nodesize=31
```

```
## [Tune-y] 11: mse.test.mean=8010794247.8236704; time: 0.0 min
## [Tune-x] 12: ntree=113; mtry=7; nodesize=33
## [Tune-y] 12: mse.test.mean=8234353188.7563391; time: 0.0 min
## [Tune-x] 13: ntree=425; mtry=4; nodesize=26
## [Tune-y] 13: mse.test.mean=8241526517.1582489; time: 0.0 min
## [Tune-x] 14: ntree=119; mtry=5; nodesize=44
## [Tune-y] 14: mse.test.mean=9155793458.8954239; time: 0.0 min
## [Tune-x] 15: ntree=474; mtry=9; nodesize=30
## [Tune-y] 15: mse.test.mean=7871056999.5831127; time: 0.0 min
## [Tune-x] 16: ntree=383; mtry=5; nodesize=45
## [Tune-y] 16: mse.test.mean=9162921917.1973724; time: 0.0 min
## [Tune-x] 17: ntree=98; mtry=9; nodesize=18
## [Tune-y] 17: mse.test.mean=7104876676.8122301; time: 0.0 min
## [Tune-x] 18: ntree=162; mtry=3; nodesize=15
## [Tune-y] 18: mse.test.mean=7963562330.2995329; time: 0.0 min
## [Tune-x] 19: ntree=339; mtry=9; nodesize=49
## [Tune-y] 19: mse.test.mean=9136753717.5987949; time: 0.0 min
## [Tune-x] 20: ntree=260; mtry=3; nodesize=24
## [Tune-y] 20: mse.test.mean=8447849156.4050350; time: 0.0 min
## [Tune-x] 21: ntree=445; mtry=9; nodesize=41
## [Tune-y] 21: mse.test.mean=8645932527.9744568; time: 0.0 min
## [Tune-x] 22: ntree=395; mtry=5; nodesize=24
## [Tune-y] 22: mse.test.mean=7823224055.2223272; time: 0.0 min
## [Tune-x] 23: ntree=294; mtry=10; nodesize=42
## [Tune-y] 23: mse.test.mean=8710023385.0719490; time: 0.0 min
## [Tune-x] 24: ntree=436; mtry=3; nodesize=36
## [Tune-y] 24: mse.test.mean=9417094187.6445255; time: 0.0 min
## [Tune-x] 25: ntree=213; mtry=9; nodesize=38
```

```
## [Tune-y] 25: mse.test.mean=8484690579.3582964; time: 0.0 min
## [Tune-x] 26: ntree=378; mtry=4; nodesize=20
## [Tune-y] 26: mse.test.mean=7771374104.1363297; time: 0.0 min
## [Tune-x] 27: ntree=365; mtry=3; nodesize=19
## [Tune-y] 27: mse.test.mean=8148368520.4292431; time: 0.0 min
## [Tune-x] 28: ntree=216; mtry=10; nodesize=30
## [Tune-y] 28: mse.test.mean=7851329397.6048946; time: 0.0 min
## [Tune-x] 29: ntree=497; mtry=8; nodesize=30
## [Tune-y] 29: mse.test.mean=7883132992.6114616; time: 0.0 min
## [Tune-x] 30: ntree=160; mtry=5; nodesize=14
## [Tune-y] 30: mse.test.mean=7162962875.7875929; time: 0.0 min
## [Tune-x] 31: ntree=374; mtry=4; nodesize=35
## [Tune-y] 31: mse.test.mean=8877329258.2792606; time: 0.0 min
## [Tune-x] 32: ntree=305; mtry=3; nodesize=22
## [Tune-y] 32: mse.test.mean=8481418415.1297550; time: 0.0 min
## [Tune-x] 33: ntree=161; mtry=9; nodesize=19
## [Tune-y] 33: mse.test.mean=7144511191.6793413; time: 0.0 min
## [Tune-x] 34: ntree=464; mtry=6; nodesize=23
## [Tune-y] 34: mse.test.mean=7657371085.8276796; time: 0.0 min
## [Tune-x] 35: ntree=102; mtry=4; nodesize=46
## [Tune-y] 35: mse.test.mean=9495121616.4848099; time: 0.0 min
## [Tune-x] 36: ntree=421; mtry=7; nodesize=34
## [Tune-y] 36: mse.test.mean=8250192634.4319096; time: 0.0 min
## [Tune-x] 37: ntree=130; mtry=6; nodesize=27
## [Tune-y] 37: mse.test.mean=7827586017.0716200; time: 0.0 min
## [Tune-x] 38: ntree=424; mtry=4; nodesize=12
## [Tune-y] 38: mse.test.mean=7166332256.5014849; time: 0.0 min
## [Tune-x] 39: ntree=385; mtry=6; nodesize=22
```



```

## [Tune-y] 39: mse.test.mean=7495472215.1052637; time: 0.0 min
## [Tune-x] 40: ntree=97; mtry=4; nodesize=31
## [Tune-y] 40: mse.test.mean=8625350472.4815712; time: 0.0 min
## [Tune-x] 41: ntree=284; mtry=3; nodesize=13
## [Tune-y] 41: mse.test.mean=7645052555.3795910; time: 0.0 min
## [Tune-x] 42: ntree=181; mtry=8; nodesize=15
## [Tune-y] 42: mse.test.mean=6727845551.9023275; time: 0.0 min
## [Tune-x] 43: ntree=487; mtry=7; nodesize=14
## [Tune-y] 43: mse.test.mean=6946887359.0446835; time: 0.0 min
## [Tune-x] 44: ntree=445; mtry=5; nodesize=33
## [Tune-y] 44: mse.test.mean=8377638581.1030493; time: 0.0 min
## [Tune-x] 45: ntree=280; mtry=10; nodesize=45
## [Tune-y] 45: mse.test.mean=8910260251.9687214; time: 0.0 min
## [Tune-x] 46: ntree=105; mtry=6; nodesize=38
## [Tune-y] 46: mse.test.mean=8745470813.6658974; time: 0.0 min
## [Tune-x] 47: ntree=458; mtry=9; nodesize=37
## [Tune-y] 47: mse.test.mean=8281123488.5573654; time: 0.0 min
## [Tune-x] 48: ntree=180; mtry=10; nodesize=29
## [Tune-y] 48: mse.test.mean=7777915296.0087576; time: 0.0 min
## [Tune-x] 49: ntree=280; mtry=4; nodesize=44
## [Tune-y] 49: mse.test.mean=9499838644.6054516; time: 0.0 min
## [Tune-x] 50: ntree=329; mtry=4; nodesize=24
## [Tune-y] 50: mse.test.mean=8161302303.4739180; time: 0.0 min
## [Tune] Result: ntree=337; mtry=9; nodesize=13 : mse.test.mean=6664777976.1
915483

sqrt(rf_tune$y)

## mse.test.mean
##      81638.09

rf.tree <- setHyperPars(rf, par.vals = rf_tune$x)

```

```

#train a model
makeatree = makeLearner("regr.randomForest", predict.type = "response")
rforest = train(rf.tree, modeling_task_train)

#make predictions
rfmodel <- predict(rforest, modeling_task_test)

#submission file
RMSE= sqrt(sum((rfmodel$data$truth-rfmodel$data$response)^2)/length(rfmodel$d
ata$response))
Rsquared = 1 - (sum((rfmodel$data$truth-rfmodel$data$response)^2)/sum((rfmodel$
data$truth-mean(rfmodel$data$response))^2))
paste0('RMSE: ',RMSE)

## [1] "RMSE: 52795.0071359748"

paste0('R.squared: ',Rsquared)

## [1] "R.squared: 0.907785473092444"

```